

Q-3 MNIST Classification

- Akshay Bankar (2019201011)

Classificaton using linear SVM

```

In [0]: 1 from sklearn import datasets, svm, metrics
        2 from sklearn.model_selection import train_test_split
        3
        4 digits = datasets.load_digits()
        5
        6 images_and_labels = list(zip(digits.images, digits.target))
        7
        8 n_samples = len(digits.images)
        9 data = digits.images.reshape((n_samples, -1))
       10
       11 classifier = svm.SVC(gamma=0.001)
       12
       13 X_train, X_test, y_train, y_test = train_test_split(
       14     data, digits.target, test_size=0.5, shuffle=False)
       15
       16 classifier.fit(X_train, y_train)
       17
       18 predicted = classifier.predict(X_test)
       19
       20 images_and_predictions = list(zip(digits.images[n_samples // 2:], predicted
       21
       22 print("Classification report for classifier %s:\n%s\n"
       23       % (classifier, metrics.classification_report(y_test, predicted)))
       24 disp = metrics.plot_confusion_matrix(classifier, X_test, y_test)
       25 disp.figure_.suptitle("Confusion Matrix")
       26 print("Confusion matrix:\n%s" % disp.confusion_matrix)
       27

```

Classification report for classifier SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=3, gamma=0.001, kernel='rbf', max_iter=-1, probability=False, random_state=None, shrinking=True, tol=0.001, verbose=False):

	precision	recall	f1-score	support
0	1.00	0.99	0.99	88
1	0.99	0.97	0.98	91
2	0.99	0.99	0.99	86
3	0.98	0.87	0.92	91
4	0.99	0.96	0.97	92
5	0.95	0.97	0.96	91
6	0.99	0.99	0.99	91
7	0.96	0.99	0.97	89
8	0.94	1.00	0.97	88
9	0.93	0.98	0.95	92
accuracy			0.97	899
macro avg	0.97	0.97	0.97	899
weighted avg	0.97	0.97	0.97	899


Confusion matrix:

```

[[87  0  0  0  1  0  0  0  0  0]
 [ 0 88  1  0  0  0  0  0  1  1]
 [ 0  0 85  1  0  0  0  0  0  0]
 [ 0  0  0 79  0  3  0  4  5  0]
 [ 0  0  0  0 88  0  0  0  0  4]
 [ 0  0  0  0  0 88  1  0  0  2]
 [ 0  1  0  0  0  0 90  0  0  0]
 [ 0  0  0  0  0  1  0 88  0  0]
 [ 0  0  0  0  0  0  0  0 88  0]
 [ 0  0  0  1  0  1  0  0  0 90]]

```

Confusion Matrix



Read binary file of the dataset

```

In [0]: 1 import time
2 stime = time.time()
3
4 import struct as st
5 import numpy as np
6 filename = {'images' : 'dataset/train-images-idx3-ubyte' , 'labels' : 'datas
7
8 labels_array = np.array([])
9
10 data_types = {
11     0x08: ('ubyte', 'B', 1),
12     0x09: ('byte', 'b', 1),
13     0x0B: ('>i2', 'h', 2),
14     0x0C: ('>i4', 'i', 4),
15     0x0D: ('>f4', 'f', 4),
16     0x0E: ('>f8', 'd', 8)}
17
18 for name in filename.keys():
19     if name == 'images':
20         imagesfile = open(filename[name], 'rb')
21     if name == 'labels':
22         labelsfile = open(filename[name], 'rb')
23
24 imagesfile.seek(0)
25 magic = st.unpack('>4B', imagesfile.read(4))
26 if(magic[0] and magic[1]) or (magic[2] not in data_types):
27     raise ValueError("File Format not correct")
28
29 nDim = magic[3]
30 print ("Data is ", nDim, "-D")
31
32 #offset = 0004 for number of images
33 #offset = 0008 for number of rows
34 #offset = 0012 for number of columns
35 #32-bit integer (32 bits = 4 bytes)
36 imagesfile.seek(4)
37 nImg = st.unpack('>I', imagesfile.read(4))[0] #num of images/labels
38 nR = st.unpack('>I', imagesfile.read(4))[0] #num of rows
39 nC = st.unpack('>I', imagesfile.read(4))[0] #num of columns
40 nBytes = nImg*nR*nC
41 labelsfile.seek(8) #Since no. of items = no. of images and is already read
42 print ("no. of images :: ", nImg)
43 print ("no. of rows :: ", nR)
44 print ("no. of columns :: ", nC)
45
46 #Read all data bytes at once and then reshape
47 images_array = 255 - np.asarray(st.unpack('>'+ 'B'*nBytes, imagesfile.read(nB
48 labels_array = np.asarray(st.unpack('>'+ 'B'*nImg, labelsfile.read(nImg))).re
49
50 print (labels_array)
51 print (labels_array.shape)
52 print (images_array.shape)

```

```

Data is 3 -D
no. of images :: 60000
no. of rows :: 28
no. of columns :: 28
[[5]
 [0]
 [4]
 ...
 [5]
 [6]
 [8]]
(60000, 1)
(60000, 28, 28)

```

Read data using MNIST library

```
In [2]: 1 from mnist import MNIST
2 import numpy as np
3 mndata = MNIST('./dataset')
4 images_array, labels_array = mndata.load_training()
5 print(np.shape(images_array))
(60000, 784)
```

Classification using CNN

Import Pytorch libraries

```
In [3]: 1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import torch.optim as optim
```

Define class for CNN

```
In [13]: 1 class Net(nn.Module):
2
3     def __init__(self):
4         super(Net, self).__init__()
5         self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
6         self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
7         self.conv2_drop = nn.Dropout2d()
8         self.fc1 = nn.Linear(320, 50)
9         self.fc2 = nn.Linear(50, 10)
10
11     def forward(self, x):
12         x = F.relu(F.max_pool2d(self.conv1(x), 2))
13         x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
14         x = x.view(-1, 320)
15         x = F.relu(self.fc1(x))
16         x = F.dropout(x, training=self.training)
17         x = self.fc2(x)
18         return F.log_softmax(x)
```

Define parameters for training

```
In [24]: 1 n_epochs = 8
2 batch_size_train = 64
3 batch_size_test = 1000
4 learning_rate = 0.001
5 momentum = 0.5
6 log_interval = 1000
```

Initialize network

```
In [14]: 1 network = Net()
          2 optimizer = optim.SGD(network.parameters(), lr=learning_rate, momentum=mome
```

Prepare train and test data

```
In [15]: 1 from sklearn.model_selection import train_test_split
          2
          3 BATCH_SIZE = 32
          4 X_train, X_test, y_train, y_test = train_test_split(np.asarray(images_array
          5
          6
          7 #####
          8 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
          9 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
         10 y_train = y_train.reshape(y_train.shape[0])
         11 y_test = y_test.reshape(y_test.shape[0])
         12 input_shape = (28, 28, 1)
         13
         14 X_train = X_train.astype('float32')
         15 X_test = X_test.astype('float32')
         16
         17 X_train /= 255
         18 X_test /= 255
         19 print('x_train shape:', X_train.shape)
         20 print('Number of images in x_train', X_train.shape[0])
         21 print('Number of images in x_test', X_test.shape[0])
         22 #####
         23
         24
         25 torch_X_train = torch.from_numpy(X_train).type(torch.FloatTensor)
         26 torch_y_train = torch.from_numpy(y_train).type(torch.LongTensor)
         27
         28 torch_X_test = torch.from_numpy(X_test).type(torch.FloatTensor)
         29 torch_y_test = torch.from_numpy(y_test).type(torch.LongTensor)
         30
         31
         32 train = torch.utils.data.TensorDataset(torch_X_train, torch_y_train)
         33 test = torch.utils.data.TensorDataset(torch_X_test, torch_y_test)
         34
         35
         36 train_loader = torch.utils.data.DataLoader(train, batch_size = BATCH_SIZE,
         37 test_loader = torch.utils.data.DataLoader(test, batch_size = BATCH_SIZE, sh
         38
```

```
x_train shape: (48000, 1, 28, 28)
Number of images in x_train 48000
Number of images in x_test 12000
```

```
In [16]: 1 train_losses = []
          2 train_counter = []
          3 test_losses = []
          4 test_counter = [i*len(images_array) for i in range(n_epochs + 1)]
```

```
In [17]: 1 def train(epoch):
2         network.train()
3         for batch_idx, (data, target) in enumerate(train_loader):
4             optimizer.zero_grad()
5             output = network(data)
6             loss = F.nll_loss(output, target)
7             loss.backward()
8             optimizer.step()
9             if batch_idx % log_interval == 0:
10                print('Train Epoch: {} [{}/{}] ({:.0f}%) \tLoss: {:.6f}'.format(
11                    epoch, batch_idx * len(data), len(train_loader.dataset),
12                    100. * batch_idx / len(train_loader), loss.item()))
13                train_losses.append(loss.item())
14                train_counter.append(
15                    (batch_idx*64) + ((epoch-1)*len(train_loader.dataset)))
```

```
In [22]: 1 def test():
2         network.eval()
3         test_loss = 0
4         correct = 0
5         with torch.no_grad():
6             for data, target in test_loader:
7                 #print(data.shape)
8                 #print(np.shape(target))
9                 #print(data)
10                output = network(data)
11                test_loss += F.nll_loss(output, target, size_average=False).item()
12                pred = output.data.max(1, keepdim=True)[1]
13                #print(pred)
14                correct += pred.eq(target.data.view_as(pred)).sum()
15            test_loss /= len(test_loader.dataset)
16            test_losses.append(test_loss)
17            print('\nTest set: Avg. loss: {:.4f}, Accuracy: {}/{} ({:.0f}%) \n'.format
18                test_loss, correct, len(test_loader.dataset),
19                100. * correct / len(test_loader.dataset)))
```

```
In [25]: 1 for epoch in range(1, n_epochs + 1):  
2         train(epoch)  
3         test()
```

/usr/lib/python3/dist-packages/ipykernel_launcher.py:18: UserWarning: Implicit dimension choice for log_softmax has been deprecated. Change the call to include dim=X as an argument.

Train Epoch: 1 [0/48000 (0%)] Loss: 0.712653

Train Epoch: 1 [32000/48000 (67%)] Loss: 0.404447

Test set: Avg. loss: 0.1710, Accuracy: 11416/12000 (95%)

Train Epoch: 2 [0/48000 (0%)] Loss: 0.818281

Train Epoch: 2 [32000/48000 (67%)] Loss: 0.493597

Test set: Avg. loss: 0.1632, Accuracy: 11432/12000 (95%)

Train Epoch: 3 [0/48000 (0%)] Loss: 0.682777

Train Epoch: 3 [32000/48000 (67%)] Loss: 0.365860

Test set: Avg. loss: 0.1592, Accuracy: 11438/12000 (95%)

Train Epoch: 4 [0/48000 (0%)] Loss: 1.049464

Train Epoch: 4 [32000/48000 (67%)] Loss: 0.418026

Test set: Avg. loss: 0.1582, Accuracy: 11428/12000 (95%)

Train Epoch: 5 [0/48000 (0%)] Loss: 0.775062

Train Epoch: 5 [32000/48000 (67%)] Loss: 0.609121

Test set: Avg. loss: 0.1514, Accuracy: 11471/12000 (96%)

Train Epoch: 6 [0/48000 (0%)] Loss: 0.769052

Train Epoch: 6 [32000/48000 (67%)] Loss: 0.510075

Test set: Avg. loss: 0.1495, Accuracy: 11476/12000 (96%)

Train Epoch: 7 [0/48000 (0%)] Loss: 0.743701

Train Epoch: 7 [32000/48000 (67%)] Loss: 0.468539

Test set: Avg. loss: 0.1452, Accuracy: 11492/12000 (96%)

Train Epoch: 8 [0/48000 (0%)] Loss: 0.347850

Train Epoch: 8 [32000/48000 (67%)] Loss: 0.464340

Test set: Avg. loss: 0.1411, Accuracy: 11506/12000 (96%)

```
In [0]: 1
```